

## ELK Installation on Centos 7

**sudo yum install vim-enhanced -y**

---

**sudo yum install java-1.8.0-openjdk**

(<https://www.digitalocean.com/community/tutorials/how-to-install-java-on-centos-and-fedora#install-openjdk-8> )

---

Nginx installation

<https://www.digitalocean.com/community/tutorials/how-to-install-nginx-on-centos-7>

---

**Refer this link for ELK installation**

<https://www.digitalocean.com/community/tutorials/how-to-install-elasticsearch-logstash-and-kibana-elastic-stack-on-centos-7>

### Step 1 — Installing and Configuring Elasticsearch

The Elastic Stack components are not available through the package manager by default, but you can install them using yum by adding Elastic's package repository.

All of the Elastic Stack's packages are signed with the Elasticsearch signing key in order to protect your system from package spoofing. Packages which have been authenticated using the key will be considered trusted by your package manager. In this step, you will import the Elasticsearch public GPG key and add the Elastic repository in order to install Elasticsearch.

Run the following command to download and install the Elasticsearch public signing key:

```
- sudo rpm --import https://artifacts.elastic.co/GPG-KEY-elasticsearch
```

Next, add the Elastic repository. Use your preferred text editor to create the file `elasticsearch.repo` in the `/etc/yum.repos.d/` directory. Here, we'll use the `vi` text editor:

```
- sudo vi /etc/yum.repos.d/elasticsearch.repo
```

To provide yum with the information it needs to download and install the components of the Elastic Stack, enter insert mode by pressing `i` and add the following lines to the file.

```
/etc/yum.repos.d/elasticsearch.repo

[elasticsearch-6.x]
name=Elasticsearch repository for 6.x packages
baseurl=https://artifacts.elastic.co/packages/6.x/yum
gpgcheck=1
gpgkey=https://artifacts.elastic.co/GPG-KEY-elasticsearch
enabled=1
autorefresh=1
type=rpm-md
```

Here you have included the human-readable name of the repo, the baseurl of the repo's data directory, and the gpg key required to verify Elastic packages.

When you're finished, press `ESC` to leave insert mode, then `:wq` and `ENTER` to save and exit the file. To learn more about the text editor `vi` and its successor `vim`, check out our [Installing and Using the Vim Text Editor on a Cloud](#) tutorial.

With the repo added, you can now install the Elastic Stack. According to the [official documentation](#), you should install Elasticsearch before the other components. Installing in this order ensures that the components each product depends on are correctly in place.

Install Elasticsearch with the following command:

```
- sudo yum install elasticsearch
```

Once Elasticsearch is finished installing, open its main configuration file, `elasticsearch.yml`, in your editor:

```
sudo vi /etc/elasticsearch/elasticsearch.yml
```

**Note:** Elasticsearch's configuration file is in YAML format, which means that indentation is very important! Be that you do not add any extra spaces as you edit this file.

Elasticsearch listens for traffic from everywhere on port 9200. You will want to restrict outside access to your Elasticsearch instance to prevent outsiders from reading your data or shutting down your Elasticsearch cluster through the REST API. Find the line that specifies `network.host`, uncomment it, and replace its value with `localhost` so it looks like this:

```
/etc/elasticsearch/elasticsearch.yml
...
network.host: localhost
...
```

Save and close `elasticsearch.yml`. Then, start the Elasticsearch service with `systemctl`:

```
- sudo systemctl start elasticsearch
```

Next, run the following command to enable Elasticsearch to start up every time your server boots:

```
- sudo systemctl enable elasticsearch
```

```
thread_pool:
bulk:
queue_size: 200
```

You can test whether your Elasticsearch service is running by sending an HTTP request:

<http://127.0.0.1:9200/> or

```
curl -X GET "localhost:9200"
```

You will see a response showing some basic information about your local node, similar to this:

## Output

```
{
  "name" : "8oSCBFJ",
  "cluster_name" : "elasticsearch",
  "cluster_uuid" : "1Nf9ZymBQaOWKpMRBfisog",
  "version" : {
    "number" : "6.5.2",
    "build_flavor" : "default",
    "build_type" : "rpm",
    "build_hash" : "9434bed",
    "build_date" : "2018-11-29T23:58:20.891072Z",
    "build_snapshot" : false,
    "lucene_version" : "7.5.0",
    "minimum_wire_compatibility_version" : "5.6.0",
    "minimum_index_compatibility_version" : "5.0.0"
  },
  "tagline" : "You Know, for Search"
}
```

Now that Elasticsearch is up and running, let's install Kibana, the next component of the Elastic Stack.

## Step 2 — Installing and Configuring the Kibana Dashboard

According to the installation order in the official documentation, you should install Kibana as the next component of the Elastic Stack. After setting Kibana up, we will be able to use its interface to search through and visualize the data stored in Elasticsearch stores.

Because you already added the Elastic repository in the previous step, you can just install the remaining components of the Elastic Stack using yum:

```
sudo yum install kibana
```

Then enable and start the Kibana service:

```
sudo systemctl enable kibana
```

**sudo systemctl start kibana**

- 

**Note:**

**(instead of localhost use 0.0.0.0 ip ES and kibana**

- **Goto /etc/elasticsearch/elasticsearch.yml. Look for value in network.host and change it to 0.0.0.0**
- **This is step if you are using Kibana. Goto /etc/kibana/kibana.yml Look for value in server.host and change it to 0.0.0.0)**

Because Kibana is configured to only listen on localhost, we must set up a reverse proxy to allow external access.

We will use Nginx for this purpose, which should already be installed on your server.

First, use the openssl command to create an administrative Kibana user which you'll use to access the Kibana web interface. As an example, we will name this account kibanaadmin, but to ensure greater security we recommend you choose a non-standard name for your user that would be difficult to guess.

The following command will create the administrative Kibana user and password, and store them in the htpasswd file. You will configure Nginx to require this username and password and read this file momentarily:

**echo "kibanaadmin:`openssl passwd -apr1`" | sudo tee -a /etc/nginx/htpasswd.users**

- 

Enter and confirm a password at the prompt. Remember or take note of this login, as you will need it to access the Kibana web interface.

===== only if you want to access via domain name start here =====

Next, we will create an Nginx server block file. As an example, we will refer to this file as example.com.conf, although you may find it helpful to give yours a more descriptive name. For instance, if you have a FQDN and DNS record set up for this server, you could name this file after your FQDN:

**sudo vi /etc/nginx/conf.d/example.com.conf**

•  
Add the following code block into the file, being sure to update example.com and www.example.com to match your server's FQDN or public IP address. This code configures Nginx to direct your server's HTTP traffic to the Kibana application, which is listening on localhost:5601. Additionally, it configures Nginx to read the htpasswd.users file and require basic authentication.

Note that if you followed the prerequisite Nginx tutorial through to the end, you may have already created this file and populated it with some content. In that case, delete all the existing content in the file before adding the following

```
example.com.conf">/etc/nginx/conf.d/example.com.conf
server {
    listen 80;

    server_name example.com www.example.com;

    auth_basic "Restricted Access";
    auth_basic_user_file /etc/nginx/htpasswd.users;

    location / {
        proxy_pass http://localhost:5601;
        proxy_http_version 1.1;
        proxy_set_header Upgrade $http_upgrade;
        proxy_set_header Connection 'upgrade';
        proxy_set_header Host $host;
        proxy_cache_bypass $http_upgrade;
    }
}
```

When you're finished, save and close the file.

Then check the configuration for syntax errors:

**- sudo nginx -t**

If any errors are reported in your output, go back and double check that the content you placed in your configuration file was added correctly. Once you see syntax is ok in the output, go ahead and restart the Nginx service:

```
- sudo systemctl restart nginx
```

By default, SELinux security policy is set to be enforced. Run the following command to allow Nginx to access the proxied service:

```
===== only if you want to access via domain name Ends here =====
```

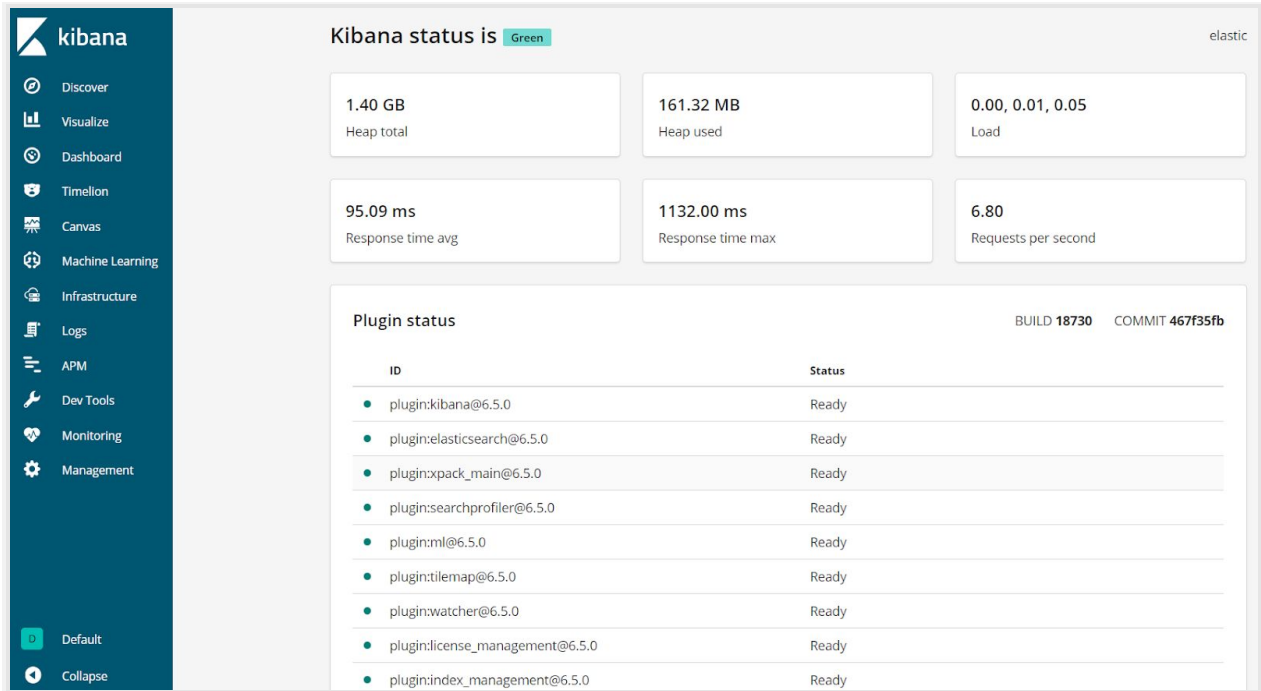
```
- sudo setsebool httpd_can_network_connect 1 -P
```

You can learn more about SELinux in the tutorial [An Introduction to SELinux on CentOS 7](#).

Kibana is now accessible via your FQDN or the public IP address of your Elastic Stack server. You can check the Kibana server's status page by navigating to the following address and entering your login credentials when prompted:

<http://localhost:5601/status>

This status page displays information about the server's resource usage and lists the installed plugins.



**Note:** As mentioned in the Prerequisites section, it is recommended that you enable SSL/TLS on your server. You can follow this tutorial now to obtain a free SSL certificate for Nginx on CentOS 7. After obtaining your SSL/TLS certificates, you can come back and complete this tutorial.

Now that the Kibana dashboard is configured, let's install the next component: Logstash.

### Step 3 — Installing and Configuring Logstash

Although it's possible for Beats to send data directly to the Elasticsearch database, we recommend using Logstash to process the data first. This will allow you to collect data from different sources, transform it into a common format, and export it to another database.

Install Logstash with this command:

- `sudo yum install logstash`



## For inserting mongo data into Elastic index using logstash (this plugin didn't work for Us)

---

<https://bqstack.com/b/detail/97/Configure-Logstash-to-send-MongoDB-data-into-Elasticsearch>

```
sudo su
cd /usr/share/logstash
bin/logstash-plugin install logstash-input-mongodb
```

We will be getting the following response:

Validating logstash-input-mongodb

Installing logstash-input-mongodb

Installation successful

It will need single config file for each collection to be inserted into elastic

*Create file in /etc/logstash/conf.d/mongodata.conf add following config*

---

```
input {
  mongodb {
    uri => 'mongodb://ntcarfte:bYzgzp9cEH@172.16.16.63:27027/ntcarfte' // mongo connection
    placeholder_db_dir => '/opt/logstash-mongodb/'
    collection => 'ntc.callpath_utilizationreport_data' // collection name
    batch_size => 500
  }
}
filter {
  mutate {
    rename => { "_id" => "mongo_id" } // _id PK of mongo collection
  }
}
output {
  elasticsearch {
    hosts => "localhost:9200"
    index => "indexname"
    doc_as_upsert => true
    document_id => "%{mongo_id}"
  }
}
```

=====

Create directory in /opt/logstash-mongodb/ and give 777 permission (delete existed .db file in that incase if you get any issue in while inserting records )

### ways to send data to elastic index

- running logstash as service

- Systemctl start logstash or
- Run config file directly to test data

/usr/share/logstash/bin/logstash -f /etc/logstash/conf.d/*mongodata.conf*

=====

### Notes

/etc/logstash/logstash.yml to debug log

**-log.level:** debug

=====

### validate all config

/usr/share/logstash/bin/logstash --path.settings /etc/logstash -f /etc/logstash/conf.d/ -t

### - open firewall ports 9200,5601

firewall-cmd --zone=public --add-port=9200/tcp --permanent

firewall-cmd --zone=public --add-port=5601/tcp --permanent

firewall-cmd --reload

iptables-save | grep 9200

### List Elastic Index

GET

127.0.0.1:9200/\_aliases/?pretty=true

## View Index data

```
get 127.0.0.1:9200/indexname/doc/_search?pretty
```

## Create Index record

```
curl -X POST '127.0.0.1:9200/ntcarfte/doc/' -H 'Content-Type: application/json' -d'
```

or

```
127.0.0.1:9200/indexname/doc/ in postman
```

```
{  
  
  "name_city_orig" : "LOSANGELES",  
  
  "future_tf_features" : null,  
  
  "date_msg" : "190711",  
  
  "record_id" : "010101",  
  
  "date" : "190712",  
  
  "@timestamp" : "2019-12-24T23:44:53.000Z",  
  
}
```

## Kibana

create stock management->index pattern in Kibana then create chart using this index

===== HOLD =====

- dynamic mapping of fields to accommodate input possible values which is skipping by logstash due to different datatypes

[https://stackoverflow.com/questions/63883823/elastic-dynamic-field-mapping/63883910?noredirect=1#comment113218646\\_63883910](https://stackoverflow.com/questions/63883823/elastic-dynamic-field-mapping/63883910?noredirect=1#comment113218646_63883910)

<https://blog.ruanbekker.com/blog/2019/04/06/elasticsearch-templates-tutorial/>

```
sudo vi /etc/elasticsearch/elasticsearch.yml
```

```
indices.memory.index_buffer_size: 60%
```

---

---

Using **monstache** instead of logstash

<https://rwyynn.github.io/monstache-site>

---

---

---

---

Install Go Lang

---

---

<https://www.digitalocean.com/community/tutorials/how-to-install-go-1-7-on-centos-7>

[https://medium.com/@nehajirafe/mongodb-to-elasticsearch-sync-using-monstache-cfe1177594b](https://medium.com/@nehajirafe/mongodb-to-elasticsearch-sync-using-monstache-cfe1177594b6)

6

go version

Step 2 : Install Monstache

- cd /usr/local/bin
- sudo wget <https://github.com/rwyynn/monstache/releases/download/v4.15.0/monstache-eee3f26.zip>
- sudo unzip monstache-eee3f26.zip
- export PATH=/usr/local/bin/build/linux-amd64:\$PATH
- cd /usr/local/bin/build/linux-amd64
- monstache -v

- =====

```
cd /usr/local/bin/build/linux-amd64
```

```
sudo vim mongo-elastic.toml
```

```
=====
mongo-url
="mongodb://admin:bYzgZp9cEH@66.160.227.205:27027/ntcarfte?authSource=
admin"
elasticsearch-url ="http://localhost:9200"
elasticsearch-max-conns = 10
resume = true
resume-name = "default"
#namespace-regex = '[ntcarfte.ntc.cdr]'
#direct-read-namespaces =
["ntcarfte.ntc.callpath_utilizationreport_data"]
direct-read-namespaces =
["ntcarfte.ntc.cdr","ntcarfte.ntc.callpath_utilizationreport_data"]
direct-read-stateful = true
#exit-after-direct-reads = true
disable-change-events = true
gzip = true
fail-fast = false
index-as-update = true
elasticsearch-retry = true
prune-invalid-json = true
index-oplog-time = true
verbose = true

[mongo-dial-settings]
timeout=15
read-timeout=7
write-timeout=7

[mongo-session-settings]
socket-timeout=30
sync-timeout=30

elasticsearch-max-bytes = 800000
elasticsearch-max-conns = 2

[gtm-settings]
```

```
channel-size = 50000
```

```
direct-read-concur = 1
direct-read-split-max = -1
direct-read-no-timeout=true
```

```
[logs]
info = "/var/log/monstache/info.log"
warn = "/var/log/monstache/warn.log"
error = "/var/log/monstache/error.log"
trace = "/var/log/monstache/trace.log"
stats = "/var/log/monstache/stats.log"
```

```
[[mapping]]
namespace = "ntcarfte.ntc.cdr"
index = "cdr_data"
type = "_doc"
```

```
[[mapping]]
namespace = "ntcarfte.ntc.callpath_utilizationreport_data"
index = "callpath_data"
type = "_doc"
```

```
=====
```

```
# Repeat mapping block it for each collection you want to
replicate
```

```
=====
```

```
Test with    => monstache -f mongo-elastic.toml &
For run it as SYSTEMD service
```

```
Create
```

```
=====
```

```
[Unit]
Description=monstache sync service
```

```
[Service]
Type=notify
```

```
ExecStart=/usr/local/bin/build/linux-amd64/monstache -f
/usr/local/bin/build/linux-amd64/mongo-elastic.toml
WatchdogSec=30s
Restart=always
```

[Install]

```
WantedBy=multi-user.target
```

```
=====
```

```
systemctl daemon-reload

systemctl enable monstache.service

systemctl start monstache.service

systemctl stop monstache.service

systemctl status monstache.service

#systemctl restart monstache.service
```

```
-----final config -----
```

```
#mongo-url="mongodb://admin:bYzgZp9cEH@66.160.227.205:27027/ntcarfte?authSource=admin"
mongo-url="mongodb://admin:wRASdYscrbj5sXSt@66.160.227.11:27018/ntcarfte?authSource=admin"
elasticsearch-url ="http://localhost:9200"
elasticsearch-max-conns = 10
resume = true
resume-name = "default"
enable-oplog = true
namespace-regex = "^ntacrftc.ntc.(cdr|callpath_utilizationreport_data)$"
direct-read-namespaces = ["ntcarfte.ntc.cdr","ntcarfte.ntc.callpath_utilizationreport_data"]
exit-after-direct-reads = true
disable-change-events = false
gzip = true
#fail-fast = false
index-as-update = false
elasticsearch-retry = true
```

```
prune-invalid-json = true
index-oplog-time = true
verbose = true
index-stats =true
stats =true
elasticsearch-max-bytes = 800000
```

```
[gtm-settings]
channel-size = 1024
direct-read-concur = 1
direct-read-split-max = -1
#direct-read-no-timeout=true
```

```
[mongo-dial-settings]
timeout=300
read-timeout=300
write-timeout=100
```

```
[mongo-session-settings]
socket-timeout=600
sync-timeout=61
```

```
[logs]
#info = "/var/log/monstache/info.log"
warn = "/var/log/monstache/warn.log"
error = "/var/log/monstache/error.log"
trace = "/var/log/monstache/trace.log"
stats = "/var/log/monstache/stats.log"
```

```
[[filter]]
#namespace = "ntcarfte.ntc.callpath_utilizationreport_data"
script = ""
module.exports = function(doc,ns) {
```

```
    var todayZero = Math.round(new Date().setHours(0,0,0,0) /1000);
```

```
    if(ns =='ntcarfte.ntc.cdr'){
        return ( ( parseInt(doc.start_epoch)) > (todayZero - 86400) );
    }
```

```
    if(ns =='ntcarfte.ntc.callpath_utilizationreport_data'){
        return (doc.inserted_at > (todayZero - 86400 ) );
    }
}
```



"""

```
[[mapping]]
namespace = "ntcarfte.ntc.cdr"
index = "cdr_data_4"
type = "_doc"
```

```
[[mapping]]
namespace = "ntcarfte.ntc.callpath_utilizationreport_data"
index = "callpath_data_4"
type = "_doc"
```

```
##### milan created this below #####
export GOBIN="$HOME/projects/bin"
export GOPATH="$HOME/projects/src"
=====
```

Incase needed for any issue

```
# source /etc/profile && source ~/.bash_profile
```

Reference

<https://github.com/rwynn/monstache-showcase/blob/master/monstache.config.toml>

=====

